

MIDS W207

Applied Machine Learning

Summer Week 07
Live Session Slides

K- Nearest Neighbors Algorithm



Setosa




Versicolor




Virginica


K- Nearest Neighbors Algorithm



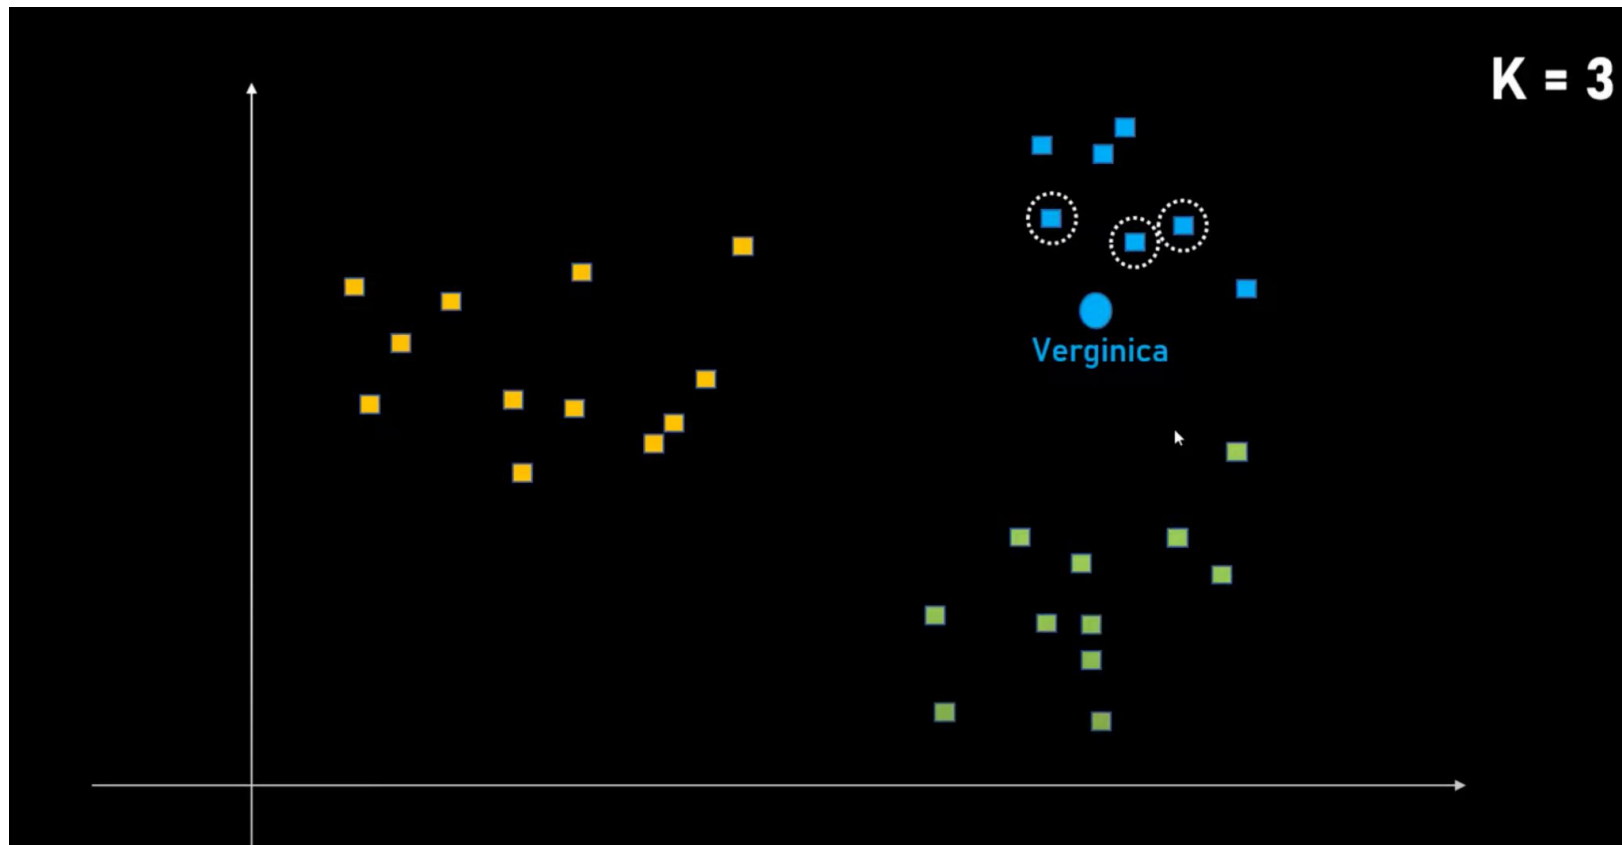
K- Nearest Neighbors Algorithm




K- Nearest Neighbors Algorithm




K- Nearest Neighbors Algorithm



K- Nearest Neighbors Algorithm

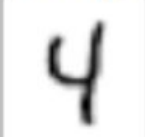
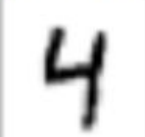
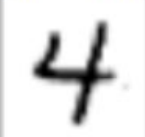


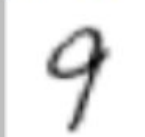


K- Nearest Neighbors Algorithm









Digit Classification







200
examples

Query	Dist: 5.47	Dist: 5.90	Dist: 5.95	Dist: 5.97	Dist: 6.32
					

1,000
examples

Query	Dist: 4.35	Dist: 4.54	Dist: 4.90	Dist: 5.19	Dist: 5.47
					

10,000
examples

Query	Dist: 3.10	Dist: 3.15	Dist: 3.95	Dist: 3.97	Dist: 3.98
					


Digit Classification

- Test set size = 10,000 digits

k = 1; Euclidean (L2) distance

Training	Error %	Time (secs)
100	30.0	0.38
1,000	12.1	2.34
10,000	5.3	28.7
60,000	2.7	2202
Deskewing	2.3	
Blurring	1.8	
Pixel shifting	1.2	

Decision Trees




Decision Tree: Customer Churn Example

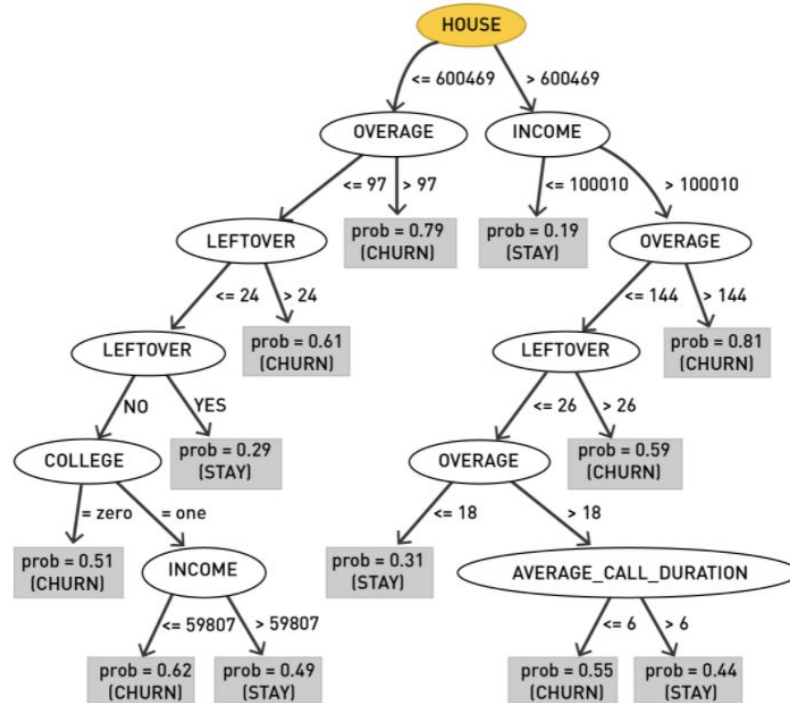
Variable	Explanation
COLLEGE	Is the customer college-educated?
INCOME	Annual income
OVERAGE	Average overcharges per month
LEFTOVER	Average number of leftover minutes per month
HOUSE	Estimated value of dwelling (from census tract)
HANDSET_PRICE	Cost of phone
LONG_CALLS_PER_MONTH	Average number of long calls (15 min or over) per month
AVERAGE_CALL_DURATION	Average duration of a call
REPORTED_SATISFACTION	Reported level of satisfaction
REPORTED_USAGE_LEVEL	Self-reported usage level
LEAVE	<i>Target variable: Did the customer stay or leave (churn)?</i>

Decision Tree: Customer Churn Example


Rank	Info. Gain	Attribute Name
1	0.0461296	HOUSE
2	0.0435518	OVERAGE
3	0.0350337	LONG CALLS PER MON
4	0.013648	LEFTOVER
5	0.0100534	INCOME
6	0.0088899	AVG CALL DURATION
7	0.007624	HANDSET PRICE
8	0.0003062	SATISFACTION
9	0.0001553	COLLEGE
10	0.0000388	USAGE LEVEL




Decision Tree: Customer Churn Example




Decision Tree: Decision Boundaries





Decision Tree: Decision Boundaries



Decision Tree: Decision Boundaries



Decision Tree: Decision Boundaries



Decision Tree: Entropy and Information Gain


Entropy is the measure of uncertainty

The more uncertainty there is about a random variable, the more information is conveyed about the value

Entropy is maximized when there is complete uncertainty through uniform distribution


0 when there is complete certainty

1



Decision Tree: Example

outlook	temperature	humidity	wind	play
sunny	hot	high	false	no
sunny	hot	high	true	no
overcast	hot	high	false	yes
rainy	mild	high	false	yes
rainy	cold	normal	false	yes
rainy	cold	normal	true	no
overcast	cold	normal	true	yes
sunny	mild	high	false	no
sunny	cold	normal	false	yes
rainy	mild	normal	false	yes
sunny	mild	normal	true	yes
overcast	mild	high	true	yes
overcast	hot	normal	false	yes
rainy	mild	high	true	no



Decision Tree: Example

The $H(\mathbf{X})$ Shannon-entropy of a discrete random variable \mathbf{X} with possible values $x_1 x_2 \dots x_n$ and probability mass function $\mathbf{P}(\mathbf{X})$ is defined as:

$$H(\mathbf{X}) = - \sum_{i=1}^n P(x_i) \log_2 P(x_i)$$

For a homogeneous dataset (all TRUE or all FALSE values, entropy is 0

For a dataset that is equally distributed (same amount of TRUEs and FALSEs, entropy is 1

A branch with entropy more than 1 needs splitting

Decision Tree: Example

outlook	temperature	humidity	wind	play
sunny	hot	high	false	no
sunny	hot	high	true	no
overcast	hot	high	false	yes
rainy	mild	high	false	yes
rainy	cold	normal	false	yes
rainy	cold	normal	true	no
overcast	cold	normal	true	yes
sunny	mild	high	false	no
sunny	cold	normal	false	yes
rainy	mild	normal	false	yes
sunny	mild	normal	true	yes
overcast	mild	high	true	yes
overcast	hot	normal	false	yes
rainy	mild	high	true	no

PLAYING GOLF

→ 9 times YES

→ 5 times NO

We just have to use the Shannon-entropy formula
to calculate the $H(x)$ values

$$H(\text{PlayingGolf}) = H(9,5) =$$

$$= -(0.64 \log_2 0.64) - (0.36 \log_2 0.36) = 0.94$$

Decision Tree: Example

outlook	temperature	humidity	wind	play
sunny	hot	high	false	no
sunny	hot	high	true	no
overcast	hot	high	false	yes
rainy	mild	high	false	yes
rainy	cold	normal	false	yes
rainy	cold	normal	true	no
overcast	cold	normal	true	yes
sunny	mild	high	false	no
sunny	cold	normal	false	yes
rainy	mild	normal	false	yes
sunny	mild	normal	true	yes
overcast	mild	high	true	yes
overcast	hot	normal	false	yes
rainy	mild	high	true	no

$$E(T,X) = \sum_x P(x) E(x)$$

We have to calculate the entropy with respect to a given predictor/feature in order to be able to calculate information gain

		PLAY GOLF	
		YES	NO
OUTLOOK	sunny	2	3
	overcast	4	0
	rainy	3	2

$$(PlayGolf, Outlook) = P(sunny)E(2,3) + P(overcast)E(4,0) + P(rainy)E(3,2)$$

$$\frac{5}{14} 0.971 + \frac{4}{14} 0 + \frac{5}{14} 0.971 = 0.6936$$

Decision Tree: Example

Information gain: the decrease in entropy after a dataset is split on an attribute/feature

→ feature/attribute with the highest information gain will be the root node in the tree

$$\begin{aligned}\text{Information Gain} &= H(\text{PlayGolf}) - E(\text{PlayGolf}, \text{Outlook}) = \\ &= 0.94 - 0.693 = 0.247\end{aligned}$$

		PLAY GOLF	
		YES	NO
OUTLOOK	sunny	2	3
	overcast	4	0
	rainy	3	2

Decision Tree: Example

Information gain: the decrease in entropy after a dataset is split on an attribute/feature

→ feature/attribute with the highest information gain will be the root node in the tree

Information Gain (outlook) = 0.247

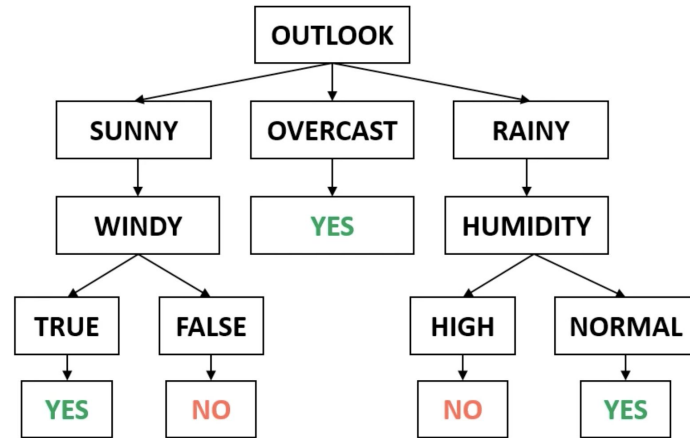
Information Gain (temperature) = 0.029

Information Gain (humidity) = 0.152


Information Gain (wind) = 0.048

Decision Tree: Example


outlook	temperature	humidity	wind	play
sunny	hot	high	false	no
sunny	hot	high	true	no
overcast	hot	high	false	yes
rainy	mild	high	false	yes
rainy	cold	normal	false	yes
rainy	cold	normal	true	no
overcast	cold	normal	true	yes
sunny	mild	high	false	no
sunny	cold	normal	false	yes
rainy	mild	normal	false	yes
sunny	mild	normal	true	yes
overcast	mild	high	true	yes
overcast	hot	normal	false	yes
rainy	mild	high	true	no




Decision Tree: Overfitting




Decision Tree: Overfitting



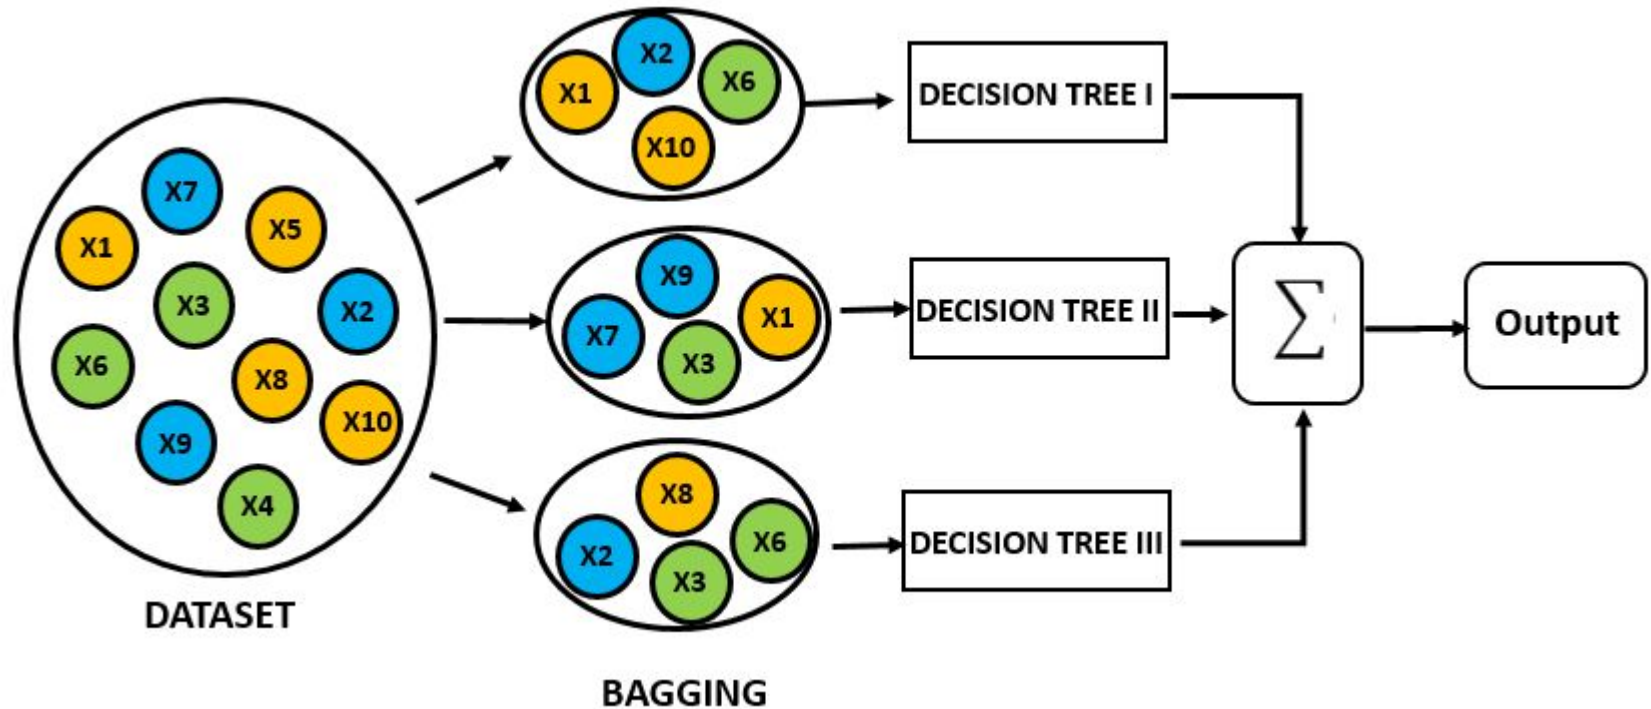
Decision Tree: Overfitting




Decision Tree: Ensemble: Bagging



Decision Tree: Ensemble: Random Forests



Decision Tree: Ensemble: Boosting



Code Review

